

R Lab 2

Zeren Li

9/2/2019

Roadmap

- tidyverse
- T-test
- Visualization using ggplot2

Advanced Data Manipulation

tidyverse

“The tidyverse is an opinionated collection of R packages designed for data science. All packages share an underlying design philosophy, grammar, and data structures.”

The core tidyverse includes the packages that you’re likely to use in everyday data analyses. As of tidyverse 1.2.0, the following packages are included in the core tidyverse:

magrittr

The magrittr package offers a set of operators which make your code more readable by:

- structuring sequences of data operations left-to-right (as opposed to from the inside and out),
- avoiding nested function calls,
- minimizing the need for local variables and function definitions, and
- making it easy to add steps anywhere in the sequence of operations.

You can think about the following sequence of actions - find key, unlock car, start car, drive to school, park.

Expressed as a set of nested functions in R pseudocode this would look like:

```
park(drive(start_car(find("keys")), to="campus"))
```

Writing it out using pipes give it a more natural (and easier to read) structure:

```
find("keys") %>%  
  start_car() %>%  
  drive(to = "campus") %>%  
  park()
```

Approaches

All of the following are fine, it comes down to personal preference:

Nested:

```
h( g( f(x), y=1), z=1 )
```

Piped:

```
f(x) %>% g(y=1) %>% h(z=1)
```

A Grammar of Data Manipulation

dplyr is based on the concepts of functions as verbs that manipulate data frames.

Single data frame functions / verbs:

- – `filter()`: filter rows by condition(s)
- – `slice()`: filter rows using index(es)
- – `select()`: select columns by name
- – `rename()`: rename variables
- – `arrange()`: reorder rows
- – `mutate()`: add new variables
- – `distinct()`: filter for unique rows
- – `sample_n()` / `sample_frac()`: randomly sample rows
- – `summarise()`: reduce variables to values
- – ... (many more)

CEO Data Analysis

Read CEO data

```
# set working directory
setwd("~/PS630-R-Lab/lab-2") # change to your own working directory
# read dta (Stata)
ceo <- read_dta("./CEOSAL2.DTA") # read CEO dataset using haven
```

`filter()` - CEOs aged over 50

```
ceo %>% filter(age > 50)
```

```
## # A tibble: 139 x 15
##   salary  age college grad comten ceoten sales profits mktval lsalary
##   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>   <dbl>
## 1    379   51      1      1      9      3    169     40    1100    5.94
## 2    651   55      1      0     22     22   1100    -54    1000    6.48
## 3   1067   64      1      1      7      7  19000    614    3900    6.97
## 4    945   59      1      0     35     10   536     24     623    6.85
## 5   1261   63      1      1     32      8   4800    191    2100    7.14
## 6   1094   64      1      1     39      5   2900    230    3900    7.00
## 7    601   54      1      1     26      7   1200     34     533    6.40
## 8    355   66      1      0     39      8    560      8     477    5.87
## 9   1200   72      1      0     37     37    796     35     678    7.09
## 10   697   51      1      0     25      1   8200    234    5700    6.55
## # ... with 129 more rows, and 5 more variables: lsales <dbl>,
```

```
## #   lmktval <dbl>, comtensq <dbl>, ceotensq <dbl>, profmarg <dbl>
```

filter() - CEOs aged over 50 and earns 1000k+

```
ceo %>% filter(age > 50, salary > 1000 )
```

```
## # A tibble: 47 x 15
##   salary  age college  grad comten ceoten sales profits mktval lsalary
##   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>   <dbl>
## 1  1067   64     1     1     7     7 19000    614   3900   6.97
## 2  1261   63     1     1    32     8  4800    191   2100   7.14
## 3  1094   64     1     1    39     5  2900    230   3900   7.00
## 4  1200   72     1     0    37    37   796     35    678   7.09
## 5  1041   63     1     1    21    11  4300     91   1400   6.95
## 6  1675   71     0     0    31    12   674    115   1200   7.42
## 7  1162   58     1     0    24     6  3800    226   1800   7.06
## 8  1627   62     1     1    13     4  8300    596   9100   7.39
## 9  1237   63     1     1    37     9  4600    108   6200   7.12
## 10 1798   66     1     1    21    14 24300    338  12500   7.49
## # ... with 37 more rows, and 5 more variables: lsales <dbl>,
## #   lmktval <dbl>, comtensq <dbl>, ceotensq <dbl>, profmarg <dbl>
```

slice() - First 10 CEOs

```
ceo %>% slice(1:10)
```

```
## # A tibble: 10 x 15
##   salary  age college  grad comten ceoten sales profits mktval lsalary
##   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>   <dbl>
## 1  1161   49     1     1     9     2  6200    966  23200   7.06
## 2   600   43     1     1    10    10   283     48   1100   6.40
## 3   379   51     1     1     9     3   169     40   1100   5.94
## 4   651   55     1     0    22    22  1100    -54   1000   6.48
## 5   497   44     1     1     8     6   351     28    387   6.21
## 6  1067   64     1     1     7     7 19000    614   3900   6.97
## 7   945   59     1     0    35    10   536     24    623   6.85
## 8  1261   63     1     1    32     8  4800    191   2100   7.14
## 9   503   47     1     1     4     4   610      7    454   6.22
## 10 1094   64     1     1    39     5  2900    230   3900   7.00
## # ... with 5 more variables: lsales <dbl>, lmktval <dbl>, comtensq <dbl>,
## #   ceotensq <dbl>, profmarg <dbl>
```

slice() - Last 5 CEOs

```
ceo %>% slice((n()-4):n())
```

```
## # A tibble: 5 x 15
##   salary  age college  grad comten ceoten sales profits mktval lsalary
##   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>   <dbl>
## 1   264   63     1     0    42     3   334     43    480   5.58
## 2   185   58     1     0    39     1   766     49    560   5.22
```

```
## 3    387    71      1    1    32    13   432     28   477    5.96
## 4   2220    63      1    1    18    18   277    -80   540    7.71
## 5    445    69      1    0    23     0   249     31   828    6.10
## # ... with 5 more variables: lsales <dbl>, lmktval <dbl>, comtensq <dbl>,
## #   ceotensq <dbl>, profmarg <dbl>
```

select() - Individual Columns

```
ceo %>% select(salary, profmarg)
```

```
## # A tibble: 177 x 2
##   salary profmarg
##   <dbl>   <dbl>
## 1  1161    15.6
## 2   600    17.0
## 3   379    23.7
## 4   651   -4.91
## 5   497     7.98
## 6  1067     3.23
## 7   945     4.48
## 8  1261     3.98
## 9   503     1.15
## 10  1094     7.93
## # ... with 167 more rows
```

select() - Exclude Columns

```
ceo %>% select(-salary, -profmarg)
```

```
## # A tibble: 177 x 13
##   age college grad comten ceoten sales profits mktval lsalary lsales
##   <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>   <dbl> <dbl>
## 1   49      1     1     9     2  6200   966  23200   7.06   8.73
## 2   43      1     1    10    10   283    48   1100    6.40   5.65
## 3   51      1     1     9     3   169    40   1100    5.94   5.13
## 4   55      1     0    22    22  1100   -54   1000    6.48   7.00
## 5   44      1     1     8     6   351    28    387    6.21   5.86
## 6   64      1     1     7     7 19000   614   3900    6.97   9.85
## 7   59      1     0    35    10   536    24    623    6.85   6.28
## 8   63      1     1    32     8  4800   191   2100    7.14   8.48
## 9   47      1     1     4     4   610     7    454    6.22   6.41
## 10  64      1     1    39     5  2900   230   3900    7.00   7.97
## # ... with 167 more rows, and 3 more variables: lmktval <dbl>,
## #   comtensq <dbl>, ceotensq <dbl>
```

```
ceo %>% select(-c(salary, profmarg))
```

```
## # A tibble: 177 x 13
##   age college grad comten ceoten sales profits mktval lsalary lsales
##   <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>   <dbl> <dbl>
## 1   49      1     1     9     2  6200   966  23200   7.06   8.73
## 2   43      1     1    10    10   283    48   1100    6.40   5.65
## 3   51      1     1     9     3   169    40   1100    5.94   5.13
```

```
## 4      55      1      0      22      22 1100      -54 1000      6.48 7.00
## 5      44      1      1       8       6  351       28  387      6.21 5.86
## 6      64      1      1       7       7 19000      614 3900      6.97 9.85
## 7      59      1      0      35      10  536       24  623      6.85 6.28
## 8      63      1      1      32       8  4800      191 2100      7.14 8.48
## 9      47      1      1       4       4   610       7   454      6.22 6.41
## 10     64      1      1      39       5  2900      230 3900      7.00 7.97
## # ... with 167 more rows, and 3 more variables: lmktval <dbl>,
## #   comtensq <dbl>, ceotensq <dbl>
```

select() - Ranges

```
ceo %>% select(salary:college)
```

```
## # A tibble: 177 x 3
##   salary age college
##   <dbl> <dbl> <dbl>
## 1  1161   49      1
## 2   600   43      1
## 3   379   51      1
## 4   651   55      1
## 5   497   44      1
## 6  1067   64      1
## 7   945   59      1
## 8  1261   63      1
## 9   503   47      1
## 10  1094   64      1
## # ... with 167 more rows
```

select() - Exclusion Ranges

```
ceo %>% select(-c(salary:college))
```

```
## # A tibble: 177 x 12
##   grad comten ceoten sales profits mktval lsalary lsales lmktval comtensq
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     1     9     2  6200   966 23200   7.06  8.73  10.1     81
## 2     1    10    10   283    48  1100   6.40  5.65   7.00    100
## 3     1     9     3   169    40  1100   5.94  5.13   7.00     81
## 4     0    22    22  1100   -54  1000   6.48  7.00   6.91    484
## 5     1     8     6   351    28   387   6.21  5.86   5.96     64
## 6     1     7     7 19000   614  3900   6.97  9.85   8.27     49
## 7     0    35    10   536    24   623   6.85  6.28   6.43   1225
## 8     1    32     8  4800   191  2100   7.14  8.48   7.65   1024
## 9     1     4     4   610     7   454   6.22  6.41   6.12     16
## 10    1    39     5  2900   230  3900   7.00  7.97   8.27   1521
## # ... with 167 more rows, and 2 more variables: ceotensq <dbl>,
## #   profmarg <dbl>
```

rename() - Change column names

```
names(ceo)

## [1] "salary"    "age"        "college"    "grad"        "comten"    "ceoten"
## [7] "sales"     "profits"    "mktval"     "lsalary"     "lsales"    "lmktval"
## [13] "comtensq"  "ceotensq"   "profmarg"

ceo_new <- ceo %>% rename(profit_margin = profmarg)
names(ceo_new)

## [1] "salary"      "age"          "college"      "grad"
## [5] "comten"      "ceoten"       "sales"        "profits"
## [9] "mktval"      "lsalary"      "lsales"       "lmktval"
## [13] "comtensq"    "ceotensq"     "profit_margin"
```

arrange() - Sort data

```
ceo %>%
  # filter if age is larger than 50
  filter(age > 50) %>%
  # sort by age and then salary
  arrange(age,salary)

## # A tibble: 139 x 15
##   salary age college grad comten ceoten sales profits mktval lsalary
##   <dbl> <dbl>   <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1    246   51     1     0     8     8     78     13     458    5.51
## 2    379   51     1     1     9     3    169     40    1100    5.94
## 3    541   51     1     0    30     4   1400     82    1200    6.29
## 4    697   51     1     0    25     1  8200    234    5700    6.55
## 5   1487   51     1     0     3     3 22200    182    2800    7.30
## 6    483   52     1     1    18    14   1000     35     548    6.18
## 7    515   52     1     1    27     1   1100     51     889    6.24
## 8    552   52     1     0    30     1  2800    308    3500    6.31
## 9    704   52     1     1     6     6    50      8     903    6.56
## 10   999   52     1     0    28    17   159     21     398    6.91
## # ... with 129 more rows, and 5 more variables: lsales <dbl>,
## #   lmktval <dbl>, comtensq <dbl>, ceotensq <dbl>, profmarg <dbl>

ceo %>%
  # filter if age is larger than 50
  filter(age > 50) %>%
  # sort by age (descend)
  arrange(desc(age),salary)

## # A tibble: 139 x 15
##   salary age college grad comten ceoten sales profits mktval lsalary
##   <dbl> <dbl>   <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1    425   86     1     1    13    13     36     11     644    6.05
## 2    396   80     1     0    58    28    513     53     963    5.98
## 3    300   77     0     0    45    26   6900    483    4700    5.70
## 4   1946   73     1     0    25    21   7800    484    8000    7.57
## 5    971   72     1     1    33    24   1400     69     609    6.88
```

```
## 6 1200 72 1 0 37 37 796 35 678 7.09
## 7 387 71 1 1 32 13 432 28 477 5.96
## 8 1675 71 0 0 31 12 674 115 1200 7.42
## 9 174 69 1 0 13 13 29 6 390 5.16
## 10 445 69 1 0 23 0 249 31 828 6.10
## # ... with 129 more rows, and 5 more variables: lsales <dbl>,
## # lmktval <dbl>, comtensq <dbl>, ceotensq <dbl>, profmarg <dbl>
```

mutate() - Modify columns

```
ceo %>%
  # add a new variable salary_l: log(salary)
  mutate(salary_l = log(salary)) %>%
  # select salary and salary_l
  select(salary, salary_l, lsalary)
```

```
## # A tibble: 177 x 3
##   salary salary_l lsalary
##   <dbl>   <dbl>   <dbl>
## 1  1161     7.06     7.06
## 2   600     6.40     6.40
## 3   379     5.94     5.94
## 4   651     6.48     6.48
## 5   497     6.21     6.21
## 6  1067     6.97     6.97
## 7   945     6.85     6.85
## 8  1261     7.14     7.14
## 9   503     6.22     6.22
## 10 1094     7.00     7.00
## # ... with 167 more rows
```

distinct() - Find unique rows

```
ceo %>%
  select(age, ceoten) %>%
  distinct() %>%
  arrange(age, ceoten)
```

```
## # A tibble: 147 x 2
##   age ceoten
##   <dbl> <dbl>
## 1  33      9
## 2  38      3
## 3  39      3
## 4  39      8
## 5  40      1
## 6  40      5
## 7  40     11
## 8  41      2
## 9  42     12
## 10 43      2
## # ... with 137 more rows
```

sample_n() sampling rows

```
ceo %>% sample_n(100)

## # A tibble: 100 x 15
##   salary age college grad comten ceoten sales profits mktval lsalary
##   <dbl> <dbl>   <dbl> <dbl>  <dbl>  <dbl> <dbl>   <dbl>   <dbl>   <dbl>
## 1    129   66     1     1     4     4    59     28    412    4.86
## 2    174   69     1     0    13    13    29      6    390    5.16
## 3   5299   64     1     0    42    13   2400    119   1500    8.58
## 4    834   58     1     0    35     1   4400     63    890    6.73
## 5    381   54     1     0    30     2   2700    386   4500    5.94
## 6    310   40     1     0    18     1   2400     60   1300    5.74
## 7   1340   55     1     0    13    10   1400    131   2900    7.20
## 8    625   57     0     0    36     9   1400     87    979    6.44
## 9    537   57     1     1    35     1  11400    210   4800    6.29
## 10   822   60     1     0    22    20    896     77    752    6.71
## # ... with 90 more rows, and 5 more variables: lsales <dbl>,
## #   lmkttval <dbl>, comtensq <dbl>, ceotensq <dbl>, profmarg <dbl>
```

summarise() - summarize data

```
ceo %>%
  # provide summary statistic: # of obs, min, max
  summarize(n = n(),
            mean = mean(salary, na.rm = T),
            min = min(salary, na.rm = T),
            max = max(salary, na.rm = T))

## # A tibble: 1 x 4
##       n mean  min  max
##   <int> <dbl> <dbl> <dbl>
## 1    177  866.   100 5299
```

Tabulate Data by grad

```
# creat your own contingency table
ceo_tab = ceo %>%
  # define subgroups
  group_by(grad) %>%
  # provide summary statistic: # of obs, min, max
  summarize(n = n(),
            mean = mean(salary, na.rm = T),
            sd = var(salary, na.rm = T) %>% sqrt(.),
            min = min(salary, na.rm = T),
            max = max(salary, na.rm = T))

ceo_tab

## # A tibble: 2 x 6
##   grad      n mean  sd  min  max
##   <dbl> <int> <dbl> <dbl> <dbl> <dbl>
## 1     0    83  868.  675.  174 5299
```



```
## 2      1      94 864. 501. 100 2265
```

Using `xtable()` to export

```
xtable(ceo_tab)
```

% latex table generated in R 3.5.1 by xtable 1.8-3 package % Wed Sep 11 04:20:58 2019

	grad	n	mean	sd	min	max
1	0.00	83	867.73	675.22	174.00	5299.00
2	1.00	94	864.21	501.39	100.00	2265.00

T-test

T-test using `t.test`

```
t.test( profmarg ~ grad, data = ceo , var.equal=TRUE, paired=FALSE)
```

```
##
## Two Sample t-test
##
## data:  profmarg by grad
## t = 0.20368, df = 175, p-value = 0.8388
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -4.774466 5.873359
## sample estimates:
## mean in group 0 mean in group 1
##      6.711907      6.162460
```

$$CI = [\bar{X} - t_{\alpha/2}S/\sqrt{n}, \bar{X} + t_{\alpha/2}S/\sqrt{n}]$$

```
# creat your own contingency table
tab = ceo %>%
  # define subgroups
  group_by(grad) %>%
  # provide summary statistic: # of obs, min, max
  summarize(n = n(),
            mean = mean(profmarg,na.rm = T),
            sd = var(profmarg,na.rm = T) %>% sqrt(.),
            min = min(profmarg,na.rm = T),
            max = max(profmarg,na.rm = T))
tab
```

```
## # A tibble: 2 x 6
##   grad      n mean    sd    min    max
##   <dbl> <int> <dbl> <dbl> <dbl> <dbl>
## 1     0    83  6.71  8.23 -48.1  23.9
## 2     1    94  6.16 23.3 -203.  47.5
```

$$t = \frac{\bar{X}_t - \bar{X}_c}{\hat{\sigma}_{\bar{X}_t - \bar{X}_c}}$$

where:

$$\hat{\sigma}_{\bar{X}_t - \bar{X}_c} = \sqrt{\frac{\hat{\sigma}_{\bar{X}_t}^2}{n_t} + \frac{\hat{\sigma}_{\bar{X}_c}^2}{n_c}}$$

```
# number of observations
n_c <- tab$n[1]
n_t <- tab$n[2]

# mean
mean_c <- tab$mean[1]
mean_t <- tab$mean[2]

# standard deviation
sd_c <- tab$sd[1]
sd_t <- tab$sd[2]

# compute sigma
sigma_tc <- sqrt(sd_c^2/n_c + sd_t^2/n_t)

# compute t-statistic
t_test <- (mean_t - mean_c)/ sigma_tc

t_test

## [1] -0.2138428
```

The degrees of freedom

R uses Welch DoF, which is estimated as follows:

$$\nu_w = \frac{\left(\frac{s_t^2}{n_t} + \frac{s_c^2}{n_c} \right)^2}{\frac{s_t^4}{n_t^2 \nu_t} + \frac{s_c^4}{n_c^2 \nu_c}}$$

```
# numerator
num = (sd_t^2/n_t + sd_c^2/n_c)^2

# denominator
den = sd_t^4/( (n_t^2) * (n_t - 1)) + sd_c^4/( (n_c^2) * (n_c - 1))

# degrees of freedom
dof = num/den
dof

## [1] 118.4336
```

P value

```
2*pt(-abs(t_test),df= dof ) # pt is the distribution function of t Distribution
## [1] 0.8310373
```

Data Visualization using ggplot2

Overview

ggplot2 is a system for declaratively creating graphics, based on The Grammar of Graphics. You provide the data, tell ggplot2 how to map variables to aesthetics, what graphical primitives to use, and it takes care of the details.

Terminology

A statistical graphic is a...

- mapping of **data**
- which may be **statistically transformed** (summarised, log-transformed, etc.)
- to **aesthetic attributes** (color, size, xy-position, etc.)
- using **geometric objects** (points, lines, bars, etc.)
- and mapped onto a specific **facet** and **coordinate system**

Ask yourself these questions before using ggplot()

- Which data is used as an input?
- Are the variables statistically transformed before plotting?
- What geometric objects are used to represent the data?
- What variables are mapped onto which aesthetic attributes?
- What type of scales are used to map data to aesthetics?

Anatomy of a ggplot


```
ggplot(  
  data = [dataframe],  
  aes(  
    x = [var_x], y = [var_y],  
    color = [var_for_color],  
    fill = [var_for_fill],  
    shape = [var_for_shape]  
  )  
) +  
  geom_[some_geom]([geom_arguments]) +  
  ... # other geometries  
  scale_[some_axis]_[some_scale]() +  
  facet_[some_facet]([formula]) +  
  ... # other options
```

Scatterplot - CEO salary and sales

Distribution

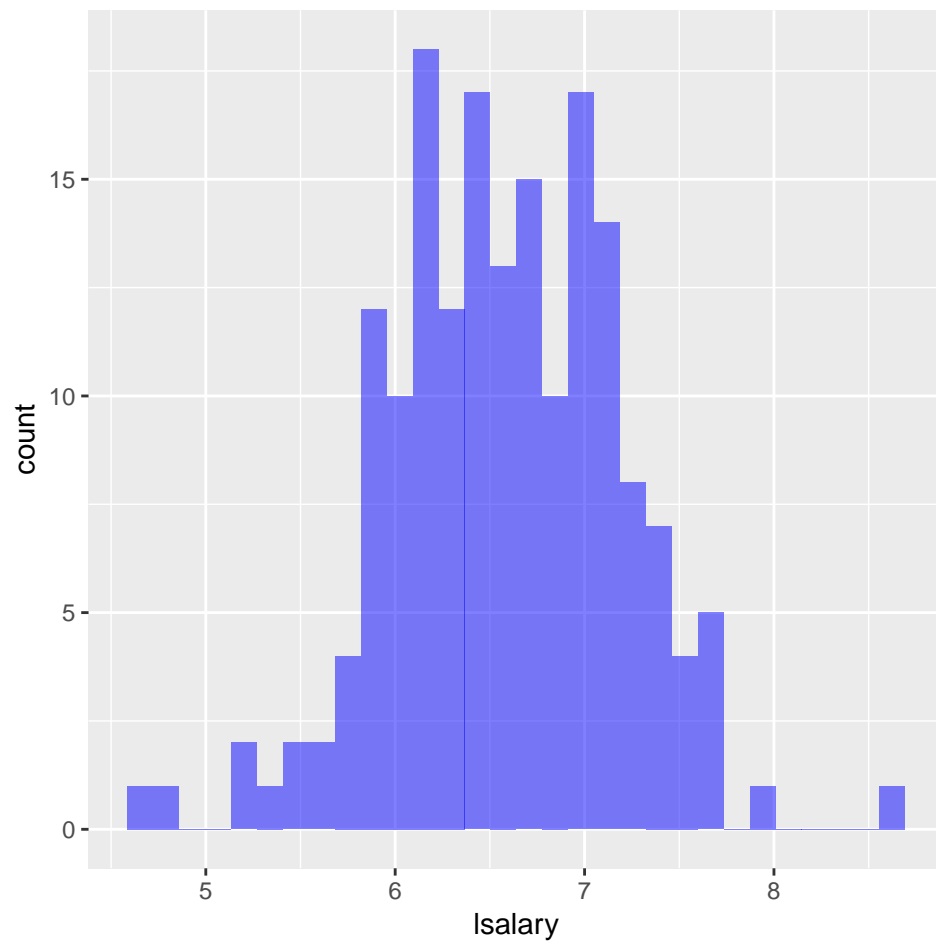
```
ggplot(data = ceo, aes(x = lsalary)) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
ggplot(data = ceo, aes(x = lsalary)) +  
  geom_histogram(alpha = .5, fill = "blue")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



Distribution

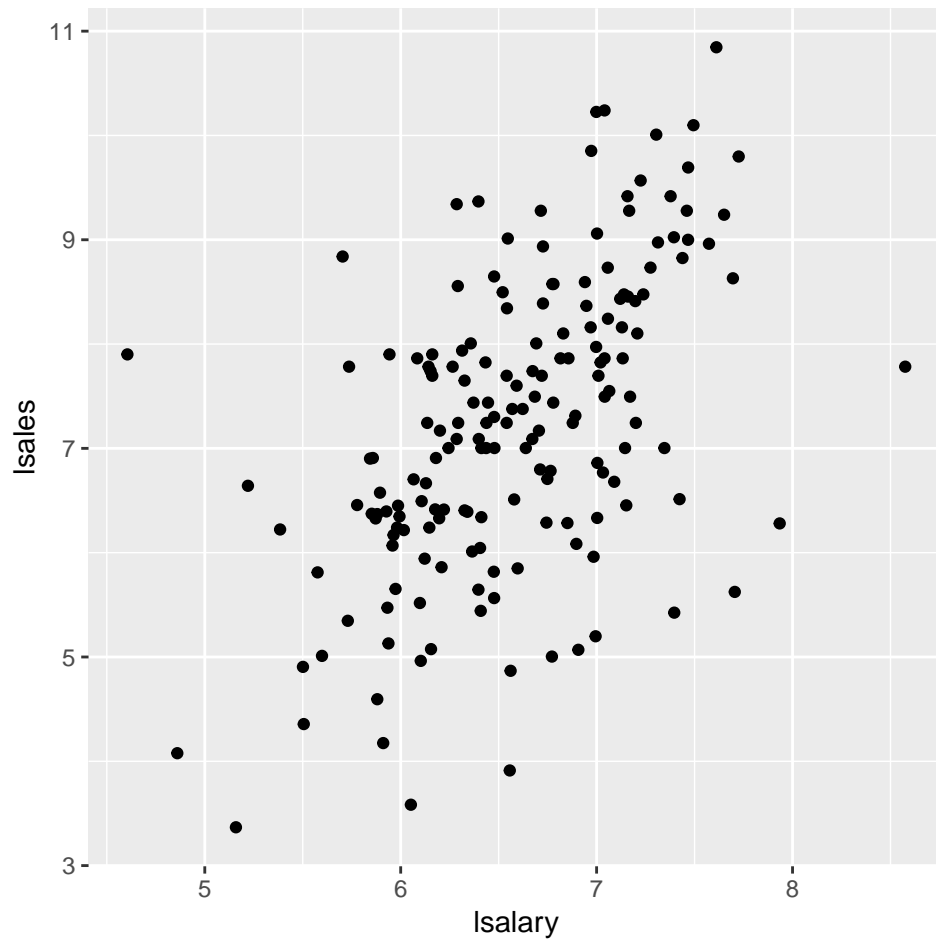
```
ggplot(data = ceo, aes(x = lsalary)) +  
  geom_density(fill = "red", alpha = .5) +  
  xlab("salary (logged)") +  
  ylab("") +  
  ggtitle("PDF of salary (logged)")
```



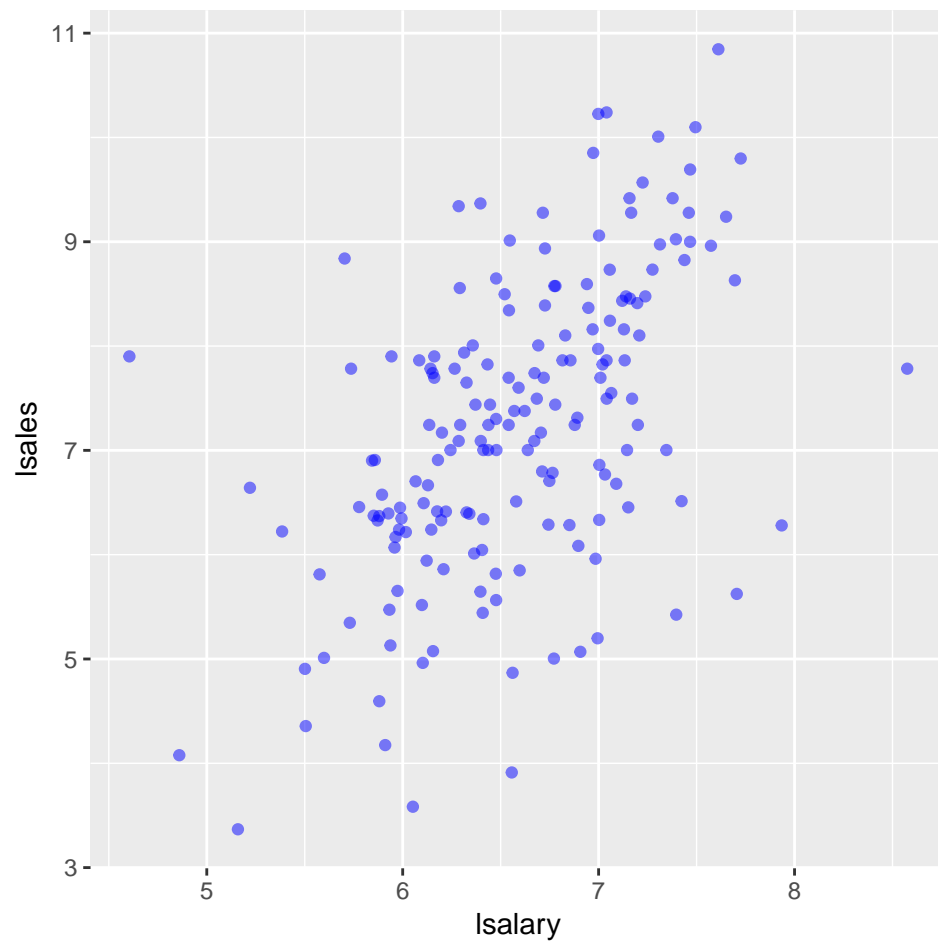
```
ggsave("./hist.pdf")
```

```
## Saving 5 x 5 in image
```

```
ggplot(data = ceo, aes(x = lsalary, y = lsales)) +  
  geom_point()
```

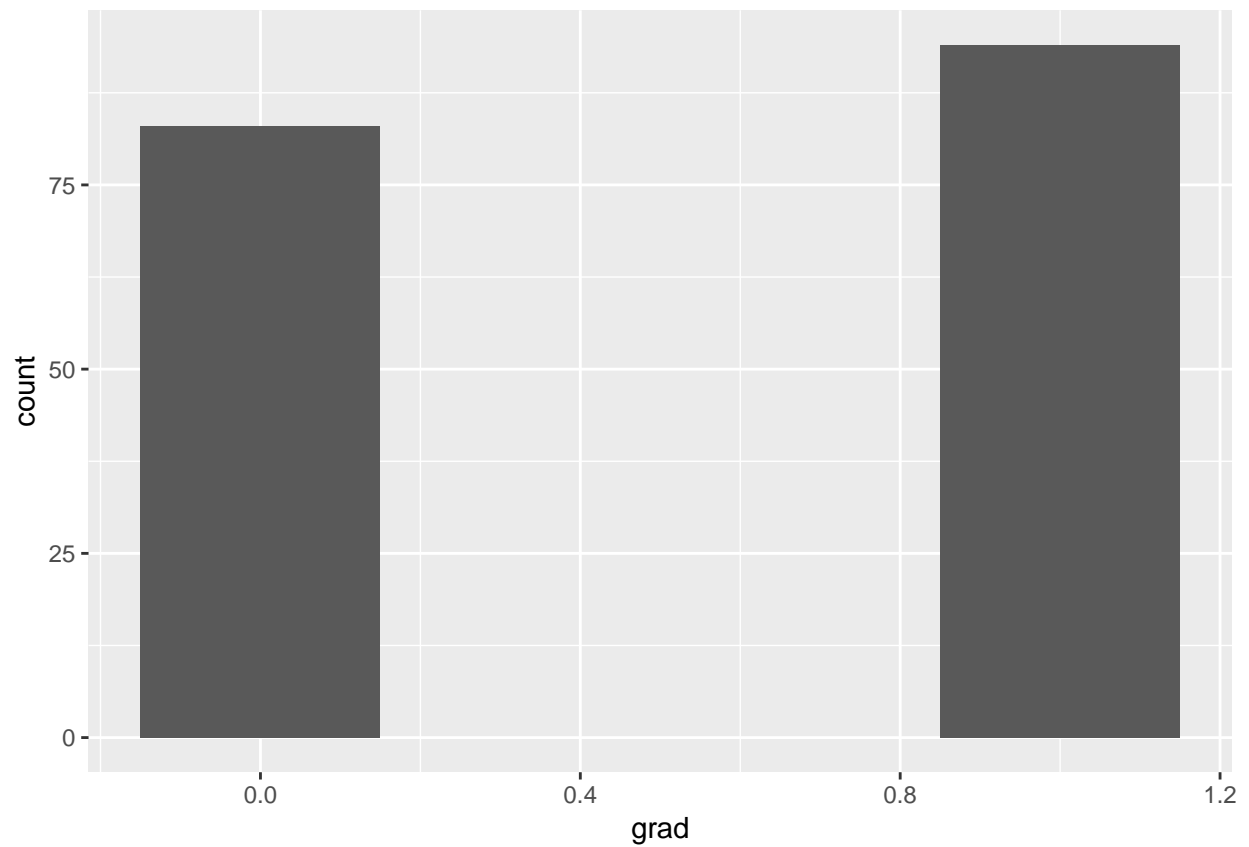


```
ggplot(data = ceo, aes(x = lsalary, y = lsales)) +  
  # specify some features  
  geom_point(alpha = 0.5, color = "blue")
```



Bar plot

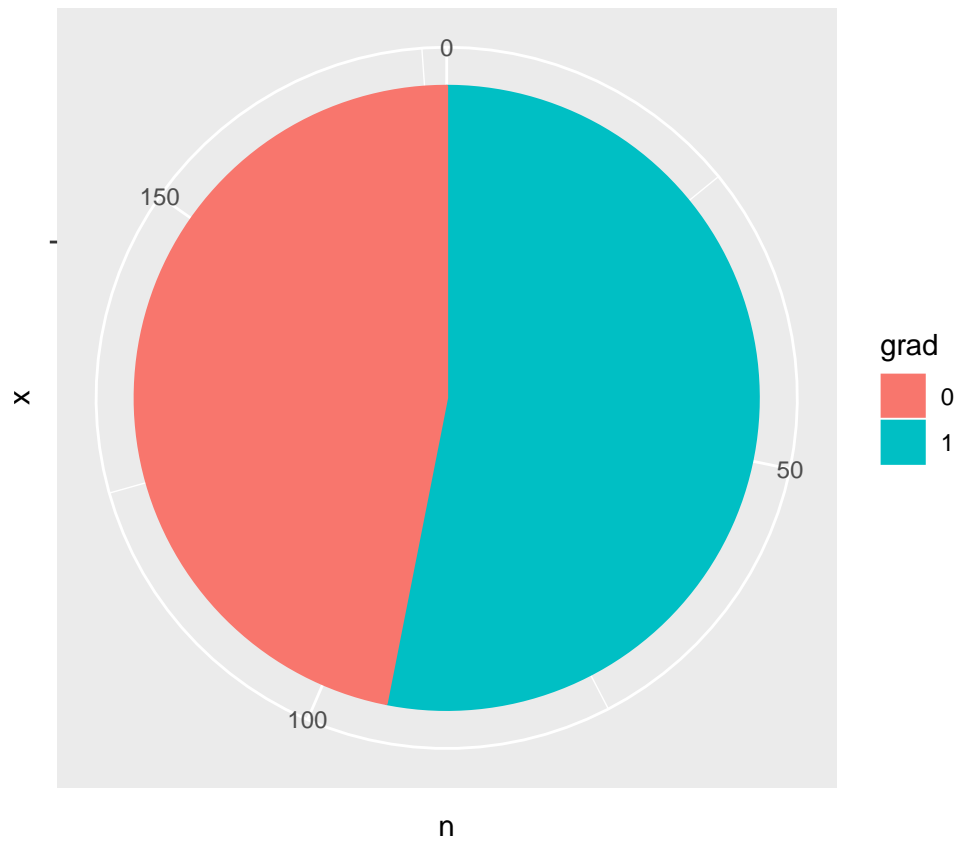
```
ggplot(data = ceo, aes(x = grad )) +  
  geom_bar(width = .3)
```

Bar plot

```
# compute mean of salary first
ceo_grad_sum <- ceo %>%
  mutate(grad = as.factor(grad)) %>%
  group_by(grad) %>%
  summarize(n = n())

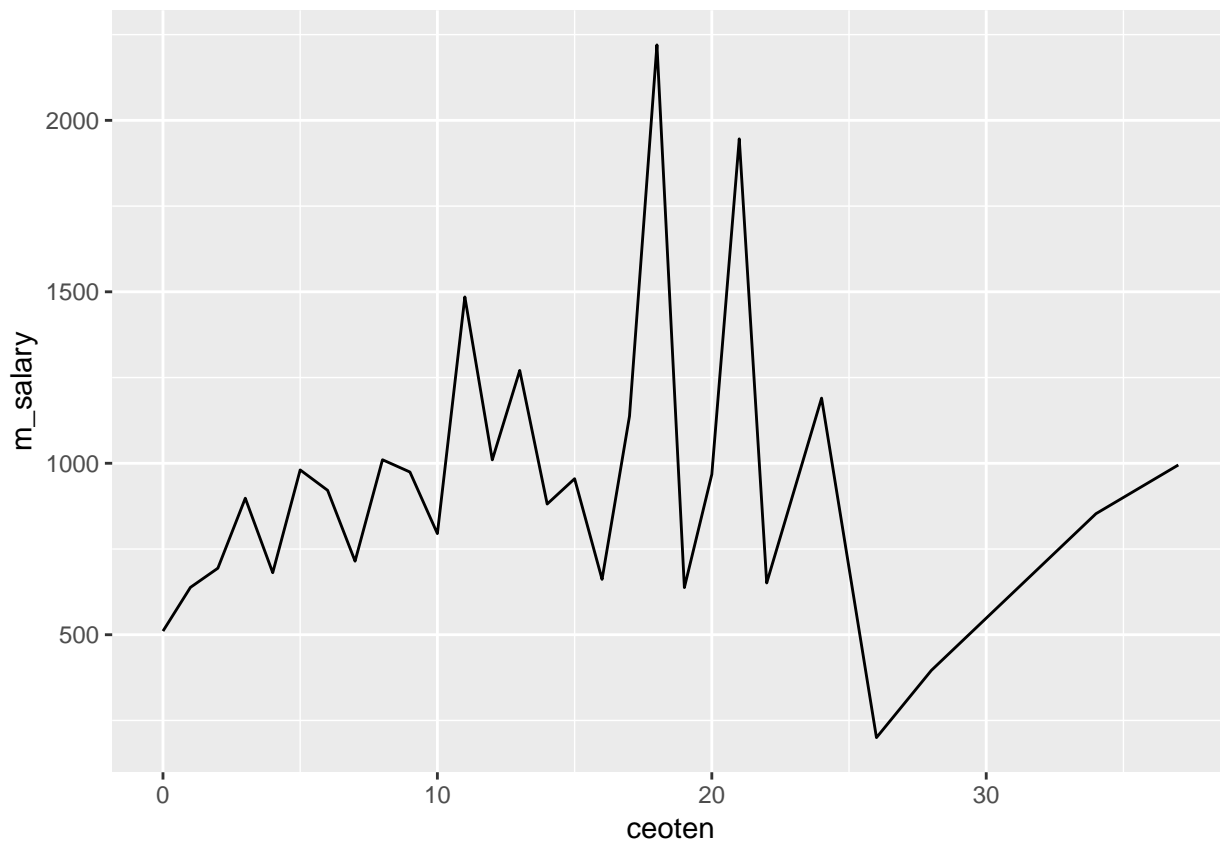
ggplot(data = ceo_grad_sum, aes(x = "", y = n, fill = grad, color = grad) ) +
  geom_bar(width = 1, stat = "identity") +
  coord_polar("y", start=0)
```



Line plot - CEO work experience and salary

```
# compute mean of salary first
ceo_sum <- ceo %>%
  group_by(ceoten) %>%
  summarise(m_salary = mean(salary, na.rm = T))

ggplot(data = ceo_sum, aes(x = ceoten, y = m_salary )) +
  geom_line()
```



A bit fancier

Resource

ggplot website: <https://ggplot2.tidyverse.org/> cheatsheet: <https://github.com/rstudio/cheatsheets/blob/master/data-visualization-2.1.pdf> top 50 visualization: <http://r-statistics.co/Top50-Ggplot2-Visualizations-MasterList-R-Code.html>